

The idiots guide to finding low-lying eigenstates of the time-independent Schrödinger equation on a non-uniform lattice, using the non-Hermitian Lanczos algorithm

Jonas Gahr Sturtzel Lunde (jonassl@berkeley.edu)

April 17, 2019

1 Intro

1.1 Diagonalization

Consider a $(N \times N)$ Hermitian and real (and therefore symmetric) matrix H , which we want to solve the eigenvalues and eigenvectors of. Since the matrix is symmetric, it is orthogonally diagonalizable, as

$$D = W^* H W \quad \Rightarrow \quad H = W D W^* \quad (1)$$

where D is a $(N \times N)$ diagonal matrix, consisting of the eigenvalues of H , while W is a orthogonal, unitary $(N \times N)$ matrix, consisting of the eigenvectors of H :

$$D = \begin{pmatrix} \lambda_1 & & & \\ & \lambda_2 & & \\ & & \ddots & \\ & & & \lambda_n \end{pmatrix} \quad W = (\mathbf{h}_1 \quad \mathbf{h}_2 \quad \cdots \quad \mathbf{h}_n)$$

1.2 Tri-diagonalization

We can also, if we wish, tridiagonalize H instead, as

$$T = V^* H V \quad \Rightarrow \quad H = V T V^* \quad (2)$$

where

$$T = \begin{pmatrix} \alpha_1 & \beta_1 & & & \\ \beta_1 & \alpha_2 & \beta_2 & & \\ & \ddots & \ddots & \ddots & \\ & & \ddots & \ddots & \ddots \\ & & & \beta_{N-1} & \alpha_N \end{pmatrix} \quad V = (\mathbf{v}_1 \quad \mathbf{v}_2 \quad \cdots \quad \mathbf{v}_n)$$

The entries of either matrix no longer has any direct meaning, but D is tridiagonal, and V is still orthogonal and unitary.

However, T and V combined, still holds access to the eigenvalues and eigenvectors of H .

If we denote $\{\mathbf{t}_i\}$ as the eigenvectors, and $\{\theta_i\}$ as the eigenvalues of T , we have that

- If θ is an eigenvalue of T , it is an eigenvalue of H .
- If \mathbf{t} is an eigenvector of T , then $\mathbf{h} = V\mathbf{t}$ is an eigenvector of H .

$$H\mathbf{h} = H V \mathbf{t} = \underbrace{(V T V^*)}_{=H} V \mathbf{t} = V T \mathbf{t} = V \theta \mathbf{t} = \theta \mathbf{h}$$

showing both that θ is an eigenvalue, and that $V\mathbf{t}$ is an eigenvector of H .

1.3

Now, consider that we instead wish to write T as a $(n \times n)$ tridiagonal matrix, where $n < N$. This will make V a $(N \times n)$ matrix. Note that V can still be orthogonal (and we will construct it as so), but it will have lost its unitary properties, such that still $V^*V = I$, but $VV^* \neq I$.

It can be shown that it is still possible to construct T as a tridiagonal matrix as

$$T = V^*HV \tag{3}$$

but note also that we no longer have $H \neq VTV^*$ due to V not being unitary anymore.

1.4 Deriving T

$$\begin{aligned} T &= V^*HV \quad \text{[Multiplying by } V^* \text{ on each side.]} \\ TV^* &= V^*HVV^* \\ TV^* &= V^*HI \end{aligned}$$

2 Stuff

2.1 Discrete Laplacian Operator

Consider the laplacian operator ∇^2 acting on some discrete system represented as a vector $\mathbf{f} = [f_0, f_1, \dots, f_N]$. Using, for instance, a 3-point central finite difference scheme, we write the laplacian acting on f_i as

$$\nabla^2 f_i = \frac{f_{i-1} - 2f_i + f_{i+1}}{h^2}$$

Let's try to make a matrix version of this operator D , such that $\frac{\partial^2 f}{\partial x^2} = D\mathbf{f}$. The i 'th output element from a matrix-vector product is the inner product between the i 'th row of the matrix, and the entire input vector. Looking at some specific output element i , we expect the picture to look something like

$$\frac{1}{h^2} D\mathbf{f} = \begin{pmatrix} \ddots & \ddots & \ddots & \ddots & \ddots \\ \ddots & 1 & -2 & 1 & \ddots \\ \ddots & \ddots & \ddots & \ddots & \ddots \end{pmatrix} \begin{pmatrix} \vdots \\ f_{i-1} \\ f_i \\ f_{i+1} \\ \vdots \end{pmatrix} = \begin{pmatrix} \vdots \\ f_{i-1} - 2f_i + f_{i+1} \\ \vdots \end{pmatrix} \quad (4)$$

The point is that the derivative of the i th element is decided by the i th row of the matrix.

Now, expanding this to the entire matrix, we get a tridiagonal matrix on the form

$$\frac{1}{h^2} D = \begin{pmatrix} -2 & 1 & 0 & 0 & \cdots \\ 1 & -2 & 1 & 0 & \cdots \\ 0 & 1 & -2 & 1 & \cdots \\ \ddots & \ddots & \ddots & \ddots & \ddots \end{pmatrix} \quad (5)$$

Generally, if an operator can be written as

$$\mathcal{O}f_i = \sum_j C_{ij} f_j$$

it has a matrix

$$O = C_{ij} = \begin{pmatrix} C_{00} & C_{01} & C_{02} & \cdots \\ C_{10} & C_{11} & C_{12} & \cdots \\ C_{20} & \ddots & \ddots & \ddots \\ \vdots & \ddots & \ddots & \ddots \end{pmatrix}$$

3 1D Radial Deuteron Potential

3.1 Open boundary conditions

Open boundary conditions works by simply considering there to be no force applied past the endpoints. This is equivalent to putting a point outside each boundary that mirrors the potential of the last poin inside our system.

4 3D Deuteron Potential

4.1 Periodic boundary conditions

Let us define $X_{x,y,z}$ or just $[x, y, z]$, for $x, y, z \in [0, N - 1]$ as our gridpoints on our 3D $N \times N \times N$ grid.

We will employ periodic boundary conditions, meaning that the ends "wrap around" each egde, such that the neighboring point of each edgepoint is the opposite edgepoint. For a grid of N points in each dimension, from 0 to $N - 1$, this can be written as

$$[-1, y, z] = [N - 1, y, z] \quad [N, y, z] = [0, y, z]$$

where the N and -1 points are artificial points, only used when considering the neighbors of the edgepoints, and not considered part of the actual system.

4.2 Unraveled indexes

Consider a 3D grid of N points in each of the 3 dimensions. The N^3 points needs to be stored in a 1D array, X_i , or simply $[i]$, with one unraveled index i . We will store them in (x,y,z) order, such that $i \in (0, N - 1)$ corresponds to the first row of x-values, at $y = 0$ and $z = 0$.

We define the conversion between $[x, y, z]$ and $[i]$ as:

$$\begin{aligned} i &= x + Ny + N^2z \\ x &= i \% N \\ y &= (i / N) \% N \\ z &= i // N^2 \end{aligned}$$

where $\%$ is the modulo operator and $//$ is an interger division.

4.3 The 7-point stencil

5 Python

5.1 Sparse Matrices

5.2 Memory Consideration

Consider a system of size $M = N^3$, where each dimension has a grid of N points. If we use a l -point Laplacian stencil, the hamiltonian will consist of $l \cdot N^3$ elements. For a 7-point stencil, and a typical $N = 100$, we get $7 \cdot 100^3 = 7 \cdot 10^6$.

The hamiltonian is stored as a sparse matrix.

The main limiting factor is going to be the storage of the V -matrix, holding the basis of the generated matrix. The V matrix has a size of $M \cdot n = N^3 \cdot n \cdot 8 \text{ B}$, which is typically of the size $100^3 \cdot 400 = 3.2 \text{ GB}$, when using 64bit floats.

6 Irregular Lattice - Laplacian

6.1 Irregular Lattice

Now, suppose you have a lattice of some dimensionality (let's say 3D). The potential we wish to study often changes rapidly in some areas, and is very uniform in others. Using a uniform lattice, we are in other words wasting lots of computational resources on un-interesting areas of the potential. We could instead let our lattice be non-uniform, with higher lattice-density in important areas.

We wish our lattice to be irregularly spaced, but constrain it to fit on some *fine lattice*, of spacing s . In other words, the spacing between each lattice point must be a multiple of s . In the simplest case, we could imagine some box around the origin where the lattice spacing is s :

$$\Delta r = 1 \cdot s, \quad -c \leq r \leq c$$

while outside the box, the lattice spacing is doubled:

$$\Delta r = 2 \cdot s, \quad |r| > c$$

6.2 Laplacian Error

The main impact on the lattice spacing is the availability of points to sample the Laplacian from. For irregular lattice-spacing, some algorithm needs to be developed to pick what points are included, and how they are weighted. For now, we will focus on an algorithm for picking the optimal spacing, with regard to the laplacian error.

Let's say you have some function on the form

$$\Psi(r) = Ae^{\kappa r} (+Be^{-\kappa r}) \quad (6)$$

for some real or complex κ . This will include both all exponential/decay functions, and all harmonic functions (and therefore most wavefunctions we are interested in). In a standard Q.M. well potential, we will get $\kappa = i\sqrt{2m|E_0|}/\hbar$ for a "trapped" solution (particle has less energy than well-depth, $E_0 < 0$), and $\kappa = \sqrt{2mE}/\hbar$ for "free" solutions (particle has more energy than well-depth $E_0 > 0$).¹

We will need to evaluate the Laplacian of this function at each latticepoint, using some finite difference scheme. Let's consider the finite difference Laplacian approximation of the 3D 7-point stencil with a uniform lattice spacing a :

$$\nabla^2 \Psi(r) = \frac{\Psi(r-a) + \Psi(r+a) - 2\Psi(r)}{a^2} + \epsilon \quad (7)$$

where ϵ is the error we have introduced with our approximation.

Now, insert the wavefunction from 6:

$$\nabla^2 \Psi(r) = \frac{A}{a^2} [e^{\kappa r} (e^{-\kappa a} + e^{i\kappa a}) - 2e^{\kappa r}] + \epsilon \quad (8)$$

Now, on the left hand side remains the analytical expression for the Laplacian. Inserting for the wavefunction, and doing a double derivation in regards to r , gives $\Psi''(r) = \kappa^2 A e^{\kappa a} = \kappa^2 \Psi(r)$. Inserting for this, and pulling $\Psi(r) = e^{\kappa r}$ out on the right side, gives

$$\kappa^2 \Psi(r) = \frac{1}{a^2} \Psi(r) [e^{-\kappa a} + e^{i\kappa a} - 2] + \epsilon \quad (9)$$

Let's Taylor-expand $e^{\pm\kappa a}$ in powers of κa , giving

$$\begin{aligned} e^{\kappa a} &= 1 + \kappa a + \frac{(\kappa a)^2}{2!} + \frac{(\kappa a)^3}{3!} + \frac{(\kappa a)^4}{4!} + \mathcal{O}(a^5) \\ e^{-\kappa a} &= 1 - \kappa a + \frac{(\kappa a)^2}{2!} - \frac{(\kappa a)^3}{3!} + \frac{(\kappa a)^4}{4!} + \mathcal{O}(a^5) \end{aligned}$$

¹Note that this is assuming the well goes to 0 at infinity. If it goes to V_0 , E_0 will instead be $E_0 = E'_0 - V_0$.

Inserting this back, we see that the first order terms cancels with the 2, and the third order term cancels with each other. We are left with

$$\kappa^2 \Psi(r) = \frac{1}{a^2} \Psi(r) \left[2 \frac{\kappa^2 a^2}{2} + 2 \frac{\kappa^4 a^4}{4!} + \mathcal{O}(a^5) \right] + \epsilon \quad (10)$$

We are left with the final error-estimate:²

$$\frac{\epsilon}{\Psi''(r)} = \frac{\kappa^2 a^2}{12} = \frac{m}{6\hbar^2} E_0 a^2 \quad \Rightarrow \quad a = \sqrt{\frac{\epsilon}{\Psi''(r)} \frac{6\hbar^2}{mE_0}} \propto \sqrt{\frac{1}{E_0}} \quad (11)$$

²(remember that $\Psi''(r) = \kappa^2 \Psi(r)$)

7 A Slightly Simplified Irregular Lattice

7.1 Point Spacing

Consider an irregularly spaced lattice, which fits on some fine grid of spacing s . We will have to decide, by some scheme, what kind of spacing $\Delta x = a \cdot s$, $a \in [1, 2, 3, \dots]$ each point on the lattice should have. One simplification of the general irregular grid, is dividing the grid into a set of b boxes, which each has a uniform grid of some lattice spacing Δx_b . This will significantly simplify calculations of the Laplacians when inside each box, as we can employ a standard, symmetric stencil. The edges will be to be handled as special cases.

7.2 Further Simplifying - Handling Edge Stencils

Yet another simplification of the box-approach can be done, which will allow us to also employ symmetric stencils at the edges. Consider an edge between two boxes, where the point-spacing changes. We could simply employ whatever stencil we were using at the lower density box onto both boxes, when handling the box edges. This will mean ignoring some close points on the higher density box, but we keep the symmetry.

Now, for this to work, we must require that all boxes have a spacing of $2^n \cdot s$, $n \in [0, 1, 2, \dots]$. If this were not true, and we considered the edge between a box with spacing of, i.e., $2s$ and $3s$, there would be no way of ignoring points in the higher density box to match the lower density one.³ While if we are looking at the edge between one $2s$ box, and one $8s$ box, we could simply ignore 3 out of every 4 points in the higher density box, and we would be left with a symmetric stencil.

7.3 Setup

One clear advantage that might seem lost when the grid was made irregular was the simplicity of neighbor-searching, due to the neighboring points in each dimension being predictably laid out in memory. In a totally chaotic grid, one would have little choice but to store points in rather arbitrary order, meaning nearest-neighbor searching would involve some form of calculating distances⁴. Another advantage of the box-implementation is that now any box will have a predictable memory layout, and neighbor-searching inside a box is simple yet again. If the second simplification mentioned above is also implemented, nearest-neighbor edge searching can also be set up in similar fashion.

8 Picking a stencil for non-rotationally symmetric points

When the points have 90 rotational symmetry, we can construct all the common stencil groups we're used to on a uniform grid. There will however be cases where no such construction is possible. It will still hold 180 degrees rotational symmetry, so we will restrict our stencils to this instead. This does however involve coming up with new, less common, sets of stencils. One important consideration that should be kept in mind in these cases is to preserve angular density of points, such that no direction holds more importance in the Laplacian.

9 Nearest Neighbor Search on a Simplified Irregular Lattice

As discussed, neighbor-searching on a simplified irregular lattice is less computationally demanding than a general case, because we can rely on the memory layout of each box. It is, however, a bit more technically challenging to implement. One demand we make, is that each point included in a stencil, must have a symmetric counterpart. We will see what implications this has.

ATM, the search goes as follows: Pick a depth D , representing the depth of points in each direction we wish to search. For instance, $D = 1$ on a regular grid would give us all points 1 step in each dimension, reducing to the 27-point stencil. Remember that we have a set of boxes, which each has a different lattice spacing a . The search then takes one of three paths:

1. We stay inside our box: If moving D steps in all directions doesn't take us to another box, finding neighbors is identical to the regular lattice case.

³We could find a lowest common denominator and ignore some points from both boxes, but that would be wasteful.

⁴Although a clever datastructure could assist in this.

2. We cross into one or several boxes with the same spacing a : If doing D steps in each direction crosses some box-borders, but they all have the same spacing, the stencil will look the same as for the regular lattice, and the implementation is not a whole lot more complicated.
3. We cross into a differently spaced box. This is the tricky part, and its explained below.

9.1 Nearest Neighbor Searches in differently spaced boxes

10 Random Notes

The code implementation for the irregular box-grid setup is such that a box always contains a con

Another thought is that the number of points which is "on the edge" of a box dramatically increases as the number of spatial dimensions increase. For instance, if the boxes have 40 gridpoints in each dimension, the outermost layer only represents $1 - 39^3/40^3 = 0.07 = 7\%$ in three spatial dimensions, while in six, they account for $1 - 39^6/40^6 = 0.14 = 14\%$. If the boxes were only of 20 points, the numbers are 14% and 26%. Edge-points might in other words be a substantial amount of the runtime.

Just confirmed that irreg and reg gives the same Hamiltonian for one big box.